

BAYESIAN GMM CLUSTERING FOR BREAST CANCER DIAGNOSIS

Giorgi Karazanashvili

ABSTRACT

This work applies a Bayesian Gaussian Mixture Model (GMM) to cluster the Wisconsin Diagnostic Breast Cancer (WDBC) dataset into malignant and benign groups without using diagnosis labels during training. A direct comparison shows that reducing dimensionality with Principal Component Analysis (PCA) before clustering degrades performance relative to clustering on the full standardized feature set. The best model, a full-covariance GMM fit on all 30 standardized features with no PCA, reaches 93.0% test accuracy, 92.9% sensitivity, 93.1% specificity, and an Adjusted Rand Index (ARI) of 0.74 against clinical labels. A Bayesian formulation with a Dirichlet prior produces an identical result, confirming the finding is not an artifact of a single optimization run.

Index Terms— Bayesian Gaussian Mixture Model, unsupervised clustering, principal component analysis, breast cancer diagnosis, Expectation-Maximization

1. INTRODUCTION

Breast cancer diagnosis from cell-nucleus morphology is traditionally framed as a supervised classification problem. This work instead asks whether an unsupervised model, given no access to diagnosis labels, can recover the malignant/benign structure directly from 30 nuclear morphology features measured from fine-needle aspirate biopsies.

The WDBC dataset contains 569 samples, each described by 10 base measurements (radius, texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry, fractal dimension), each summarized by its mean, standard error, and “worst” (mean of the three largest) value, for 30 features in total.

A common assumption in high-dimensional clustering is that reducing dimensionality with PCA before fitting a mixture model improves results, since it removes redundant, correlated features. This work tests that assumption directly and finds it does not hold for this dataset: PCA consistently reduces clustering quality relative to using the full feature space.

2. METHODOLOGY

2.1. Data and preprocessing

Samples were split 80/20 into train and test sets, stratified by diagnosis. All features were standardized (zero mean, unit variance) using statistics from the training set only.




Fig. 1. Class distribution: 357 benign, 212 malignant samples.




Fig. 2. Correlation matrix across all 30 features. Substantial redundancy (e.g. radius, perimeter, and area all measure size) motivates testing PCA as a preprocessing step.

2.2. PCA branch

A PCA projection retaining 95% of variance was fit on the training set, reducing the feature space from 30 to 10 dimensions. Figure 3 shows the cumulative variance explained.




Fig. 3. Variance explained per principal component (left) and cumulative variance with the 95% threshold (right).

Projecting onto the first two principal components alone shows visually separable clusters (Figure 4), which initially suggested PCA would help clustering. The experimental

results in Section 3 show this visual separability does not translate into better clustering performance.




Fig. 4. 2D PCA projection colored by true diagnosis. Classes appear visually separable in just two components.

2.3. Model

A two-component Gaussian Mixture Model was fit via Expectation-Maximization, with covariance type (full vs. diagonal) and preprocessing (raw 30D features vs. PCA at various dimensionalities) varied as ablations. Cluster identity is arbitrary in unsupervised fitting, so predicted cluster labels were aligned to true diagnosis labels by choosing whichever assignment maximizes agreement before computing accuracy or the confusion matrix. A Bayesian variant (BayesianGaussianMixture with a Dirichlet prior) was also fit on the best-performing configuration to confirm the result holds under a Bayesian treatment.

3. EXPERIMENTAL RESULTS

3.1. Effect of PCA and covariance structure

Five configurations were compared on the held-out test set: K-means and GMM on the PCA-reduced (10D) features, GMM restricted to just the first two principal components, GMM with diagonal covariance on PCA features, and GMM with diagonal and full covariance directly on the raw standardized 30D features.




Fig. 5. Adjusted Rand Index across preprocessing and covariance configurations. Full-covariance GMM without PCA (green) matches the

best PCA-based result while diagonal covariance on PCA features (red) performs worst.

Full-covariance GMM on the raw 30D features achieves the same ARI (0.74) as K-means on PCA features, without discarding any dimensions. Restricting to PCA features with diagonal covariance performs worst of all configurations (ARI 0.10), despite operating on the same 2D projection that looked visually separable in Figure 4. This indicates that PCA's top components, chosen to capture total variance, are not aligned with the directions that separate malignant from benign samples; collapsing to fewer components discards class-discriminative information that full covariance on the original feature space preserves.

3.2. Final model performance

The final model, full-covariance GMM on all 30 standardized features, no PCA, was evaluated on the test set.




Fig. 6. Test set confusion matrix for the final model. 93.0% accuracy, 92.9% sensitivity, 93.1% specificity.

Sensitivity (malignant recall) of 92.9% and specificity (benign recall) of 93.1% show the model is not biased toward either class, despite the moderate class imbalance in the original dataset (62.7% benign). A Bayesian GMM with a Dirichlet prior, fit on the same 30D features, produced an identical test ARI (0.74) and accuracy (93.0%), confirming the result is robust to the Bayesian formulation.




Fig. 7. True labels (left), final model cluster assignments (center), and misclassified test points (right), shown in 2D PCA space for visualization only, PCA was not used in fitting this model.

4. DISCUSSION

4.1. *Why PCA underperforms here*

PCA selects directions of maximum total variance, which need not align with directions that separate classes. In the WDBC feature space, the leading principal components capture overall tumor size and shape variation, but discriminative information distinguishing malignant from benign tissue is distributed across additional directions that PCA truncates. Full covariance on the raw feature space can model correlations between all 30 original features directly, recovering signal that a 10D or 2D PCA projection discards.

4.2. *Unsupervised performance*

The final model reaches 93.0% accuracy without ever observing a diagnosis label during training. This is comparable to many supervised baselines on this dataset, and demonstrates that the malignant/benign distinction is largely recoverable from unsupervised structure in cell-nucleus morphology alone.

4.3. *Limitations and future work*

The two-component assumption matches the clinical binary diagnosis, but was not selected via formal model selection criteria such as BIC or AIC across a range of K . The Dirichlet Process variant of the Bayesian GMM, which infers the number of effective clusters automatically, did not improve on the fixed $K=2$ result in preliminary experiments and was not pursued further here. Future work could apply formal model-selection criteria across covariance structures and component counts, and test whether supervised feature selection (rather than unsupervised PCA) better preserves class-discriminative information under dimensionality reduction.

REFERENCES

- [1] W. N. Street, W. H. Wolberg, and O. L. Mangasarian, "Nuclear feature extraction for breast tumor diagnosis," in *Proc. IS&T/SPIE*, vol. 1905, pp. 861–870, 1993.
- [2] C. Fraley and A. E. Raftery, "Model-based clustering, discriminant analysis, and density estimation," *J. Amer. Statist. Assoc.*, vol. 97, no. 458, pp. 611–631, 2002.
- [3] D. M. Blei and M. I. Jordan, "Variational inference for Dirichlet process mixtures," *Bayesian Analysis*, vol. 1, no. 1, pp. 121–143, 2006.
- [4] D. Dua and C. Graff, "UCI Machine Learning Repository: Breast Cancer Wisconsin (Diagnostic) Dataset," University of California, Irvine, 2019. [Online]. Available: [https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+\(Diagnostic\)](https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+(Diagnostic)).
- [5] F. Pedregosa et al., "Scikit-learn: Machine Learning in Python," *J. Mach. Learn. Res.*, vol. 12, pp. 2825–2830, 2011.